



浙江大学爱丁堡大学联合学院

ZJU-UoE Institute

ADS2 Lecture 2.1

Comparing multiple means

Dr Duncan MacGregor duncan.macgregor@ed.ac.uk

Semester 2, Week 1

2023-24

So, you know how to do a t-test



So, you know how to do a t-test



But what if you want to compare more than two means?

Learning Objectives

After this lecture you should be able to ...

- Design and interpret a simulation-based hypothesis test
- Use a simulation-based test to compare more than two means
- Discuss limitations of t-tests
- Discuss problems around multiple testing

What if I want to compare more than two groups?

Can you name situations where this problem would arise?

What if I want to compare more than two groups?

Can you name situations where this problem would arise?

More than 2 populations

- Effect of different drugs on recovery from injury

What if I want to compare more than two groups?

Can you name situations where this problem would arise?

More than 2 populations

- Effect of different drugs on recovery from injury
- Feeding behaviour of different bird species
- Comparing healthcare in China, the UK and Austria
- ...

What if I want to compare more than two groups?

Can you name situations where this problem would arise?

More than 2 populations

- Effect of different drugs on recovery from injury
- Feeding behaviour of different bird species
- Comparing healthcare in China, the UK and Austria
- ...

More than 1 predictive variable

- Effect of diet *and* exercise on health

What if I want to compare more than two groups?

Can you name situations where this problem would arise?

More than 2 populations

- Effect of different drugs on recovery from injury
- Feeding behaviour of different bird species
- Comparing healthcare in China, the UK and Austria
- ...

More than 1 predictive variable

- Effect of diet *and* exercise on health
- Effect of genetic background *and* drugs on stress levels
- Differences in height by gender and birth province
- ...

OK, we can't just run a t-test here


Why not?

OK, we can't just run a t-test here


Why not?

OK. But maybe we can run several t-tests?

Example: Comparing four groups



Example: Comparing four groups



How many t-tests would you need to run?

Example: Comparing four groups

Comparing:

A to B

A to C

A to D

B to C

B to D

C to D

t-test review

What is the probability of getting a false-positive result if there really is no difference?

If you are not sure, think about what happens when we do a t-test. What does your p-value mean? How do you use it to decide?



t-test review

- When computing the p-value we ask:
If H_0 is true, what is the probability of seeing a result as or more extreme as the one we saw in our experiment?

t-test review

- When computing the p-value we ask:
If H_0 is true, what is the probability of seeing a result as or more extreme as the one we saw in our experiment?
- We compare the p-value to the “significance level” α to decide on whether or not to reject H_0

t-test review

- When computing the p-value we ask:
If H_0 is true, what is the probability of seeing a result as or more extreme as the one we saw in our experiment?
- We compare the p-value to the “significance level” α to decide on whether or not to reject H_0
- If we set α to 0.05, this means that we are happy with p-values lower than 0.05


t-test review

- When computing the p-value we ask:
If H_0 is true, what is the probability of seeing a result as or more extreme as the one we saw in our experiment?
- We compare the p-value to the “significance level” α to decide on whether or not to reject H_0
- If we set α to 0.05, this means that we are happy with p-values lower than 0.05
- But every so often, even if H_0 is true, we will see a result as or more extreme as the one we saw. How often exactly?

t-test review

- When computing the p-value we ask:
If H_0 is true, what is the probability of seeing a result as or more extreme as the one we saw in our experiment?
- We compare the p-value to the “significance level” α to decide on whether or not to reject H_0
- If we set α to 0.05, this means that we are happy with p-values lower than 0.05
- But every so often, even if H_0 is true, we will see a result as or more extreme as the one we saw. How often exactly?
- By setting α , we accept a certain risk of seeing a false positive (given H_0 is true). This is exactly what α is.

Assume there really is no difference



In 6 t-tests with $\alpha = 0.05$ for each, what is the probability of getting at least one false positive result?

Let's do the maths

$P(\text{at least one false positive})$

$= 1 - P(\text{no false positives})$

$= 1 - (0.95)^6$


$= 1 - 0.735$

$= 26.5 \%$

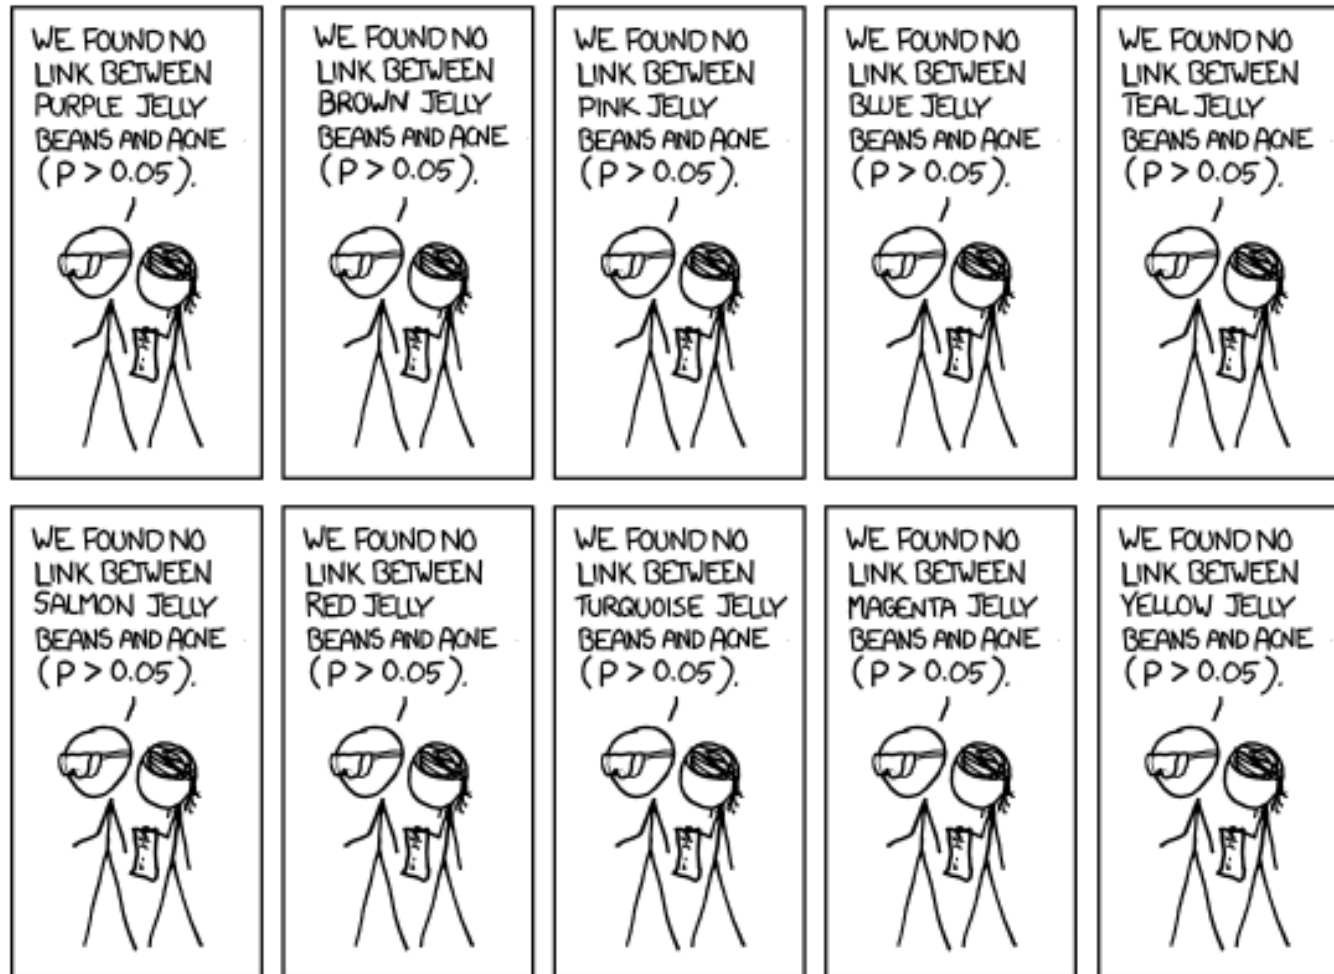
Very relevant xkcd comic

SIGNIFICANT

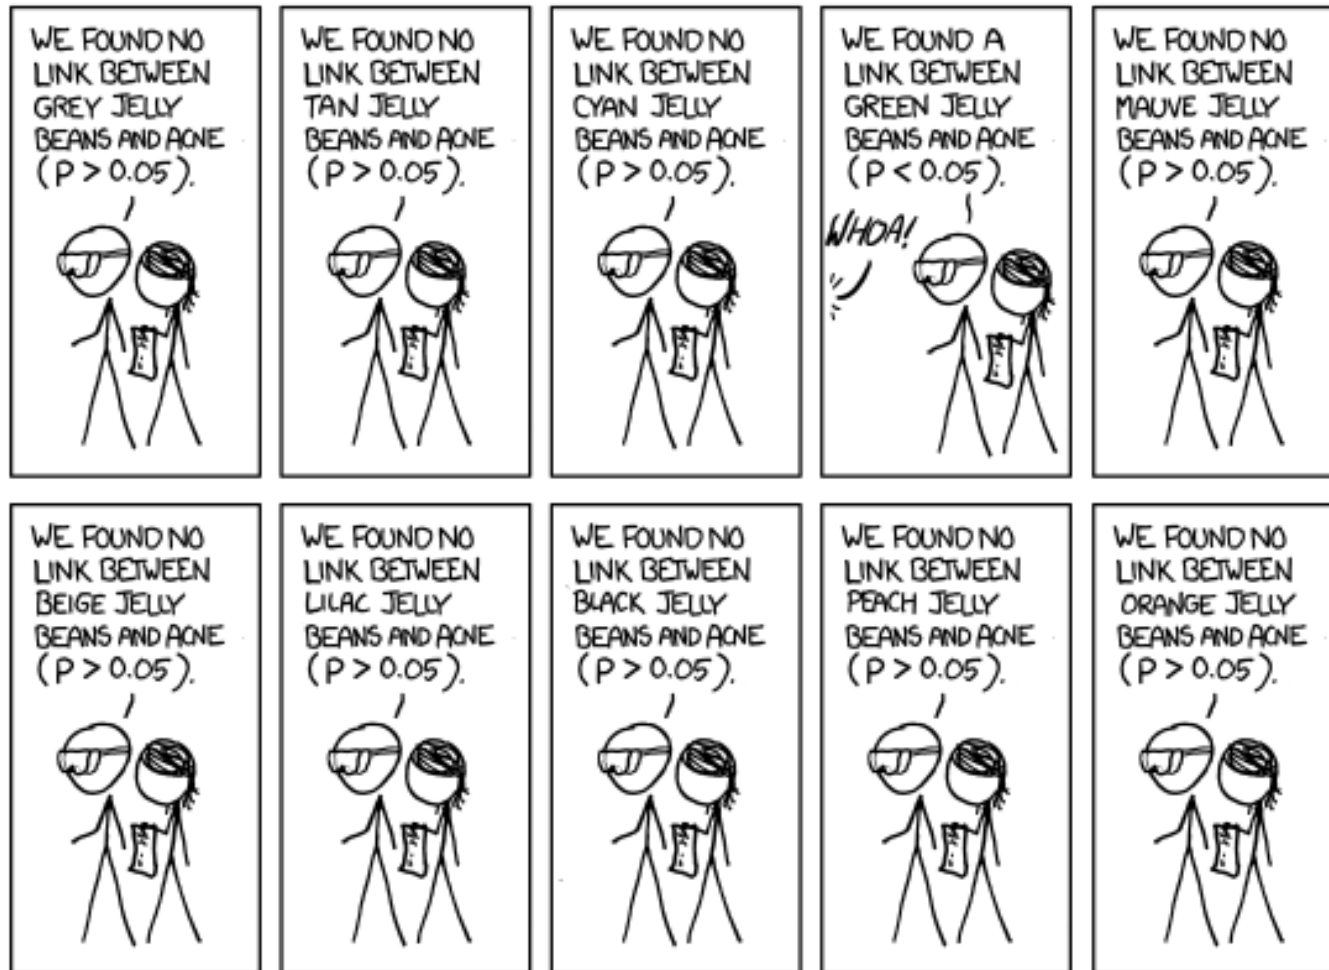
|< < PREV RANDOM NEXT > |




Very relevant xkcd comic



Very relevant xkcd comic



Very relevant xkcd comic



Yes, OK, but . . .

If we can't do a bunch of t-tests, what other option do we have?

Key idea

Looking not at group means, but at variation between individuals.

Key idea

Looking not at group means, but at variation between individuals.

Key question

If I select two individuals from different groups, are they going to be more different than if I select two individuals from the same group?

Can you think of another way of phrasing this question?

Key idea

Looking not at group means, but at variation between individuals.

Key question


If I select two individuals from different groups, are they going to be more different than if I select two individuals from the same group?

Can you think of another way of phrasing this question?

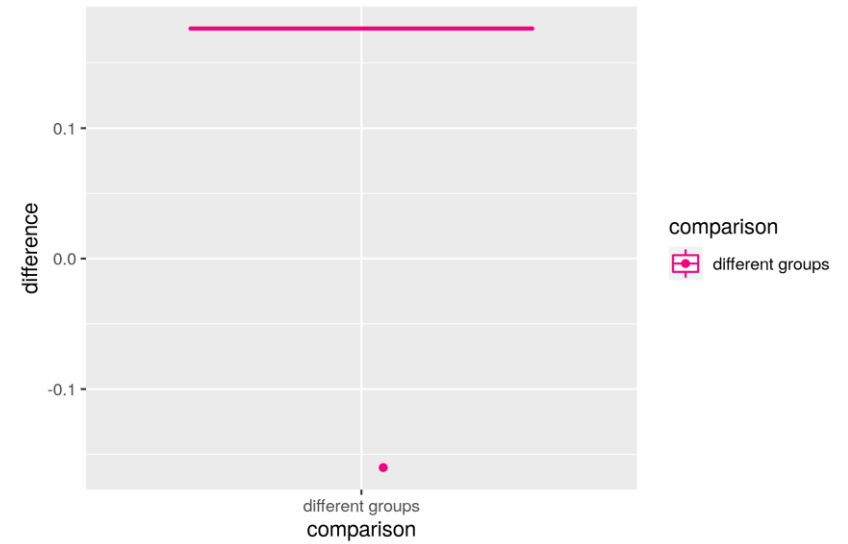

Alternative formulation

How much of the variation between individuals is explained by differences **between** groups (as opposed to differences **within** the same group)?



Looking at differences within and between groups





Looking at differences within and between groups



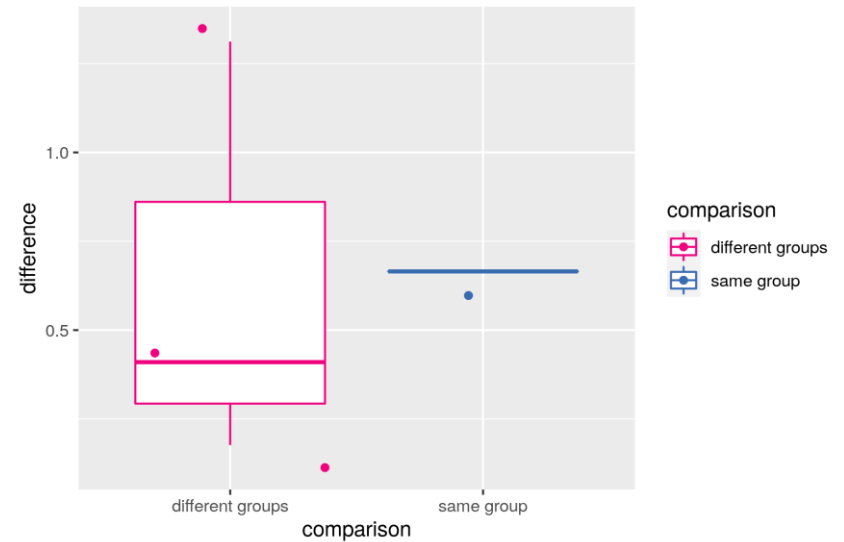

Looking at differences within and between groups



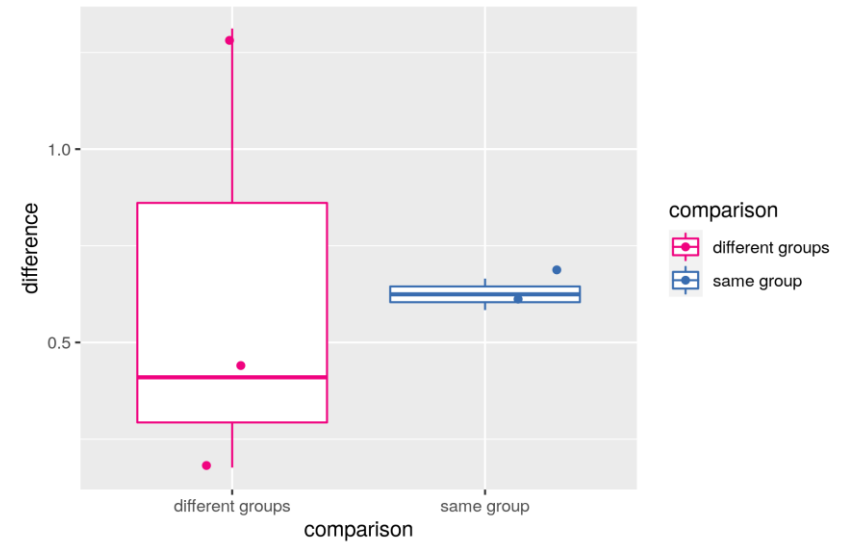

Looking at differences within and between groups





Looking at differences within and between groups





Looking at differences within and between groups





Looking at differences within and between groups



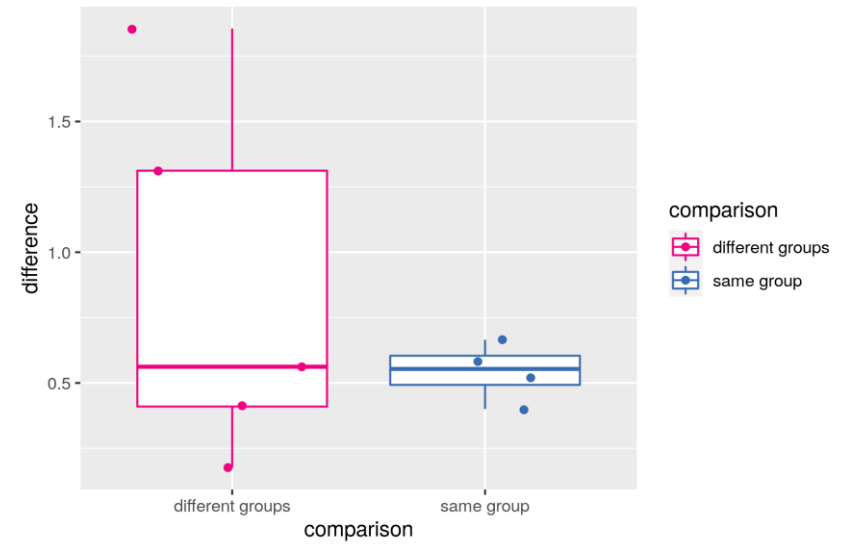

Looking at differences within and between groups





Looking at differences within and between groups




Looking at differences within and between groups




Looking at differences within and between groups



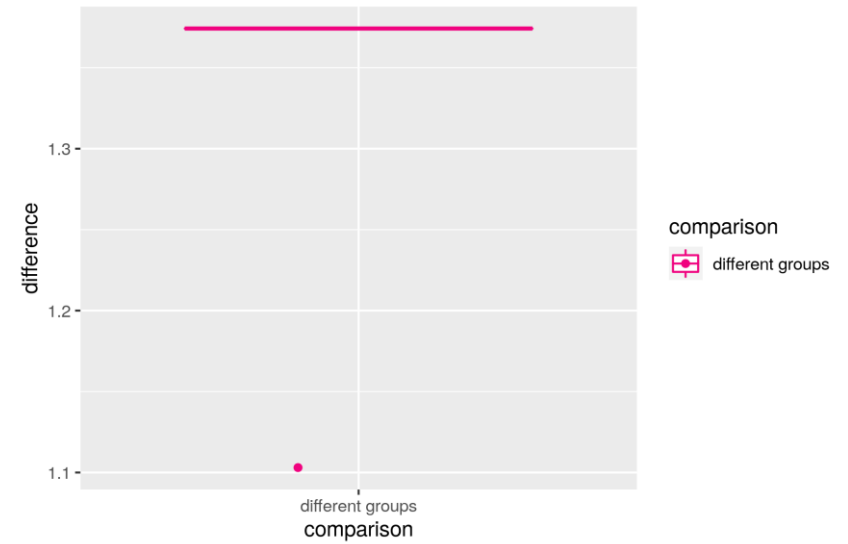

After about 1000 iterations ...





Looking at differences within and between groups - example 2





Looking at differences within and between groups





Looking at differences within and between groups





Looking at differences within and between groups





Looking at differences within and between groups





Looking at differences within and between groups





Looking at differences within and between groups





Looking at differences within and between groups



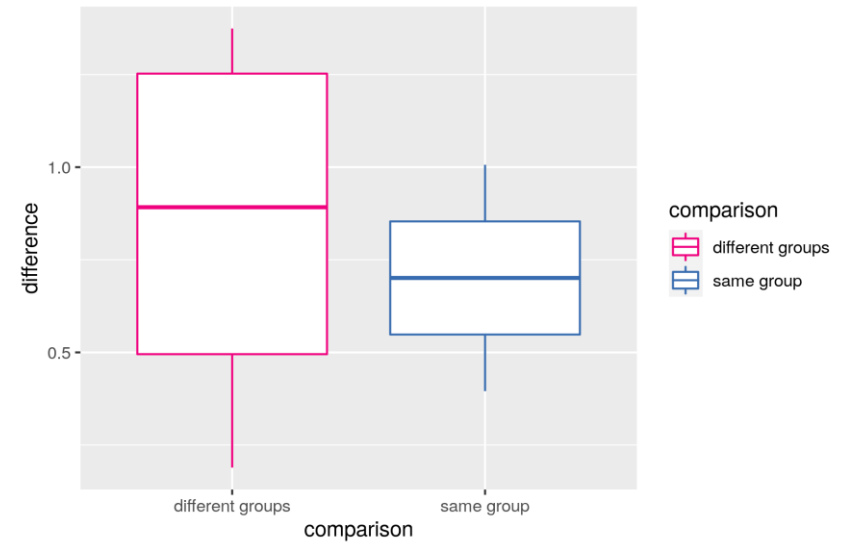

Looking at differences within and between groups




Looking at differences within and between groups



Looking at differences within and between groups




After about 1000 iterations ...



How does this help?


Why is it useful to look at within-group and between group differences?



It once more gives us two things to compare, reducing the problem to a single test.


Preview: This week's practical

Using a simulation-based approach to determine whether there are differences between 3 groups.



Preview: Next week's lecture

A more formal look at ANalysis Of VAriance



What questions do you have?

After this week you should be able to ...

- Design and interpret a simulation-based hypothesis test
- Use a simulation-based test to compare more than two means
- Discuss limitations of t-tests
- Discuss problems around multiple testing

Acknowledgements and Image credits

This lecture uses materials from ADS2 lectures by Melanie Stefan. Where not otherwise indicated, images are also from those lectures.

- ANOVA. From a lecture by Nicola Romano for ABMS2, 2018/10.
- Four glasses of beer. By FASTILY - Own work, CC BY-SA 4.0,
<https://commons.wikimedia.org/w/index.php?curid=78103101>
- R plots of various data sets. My own work, 2019-2020. CC BY-SA 3.0.
- Significant. By Randall Munroe, xkcd comics <https://xkcd.com/882/>. This work is licensed under a Creative Commons Attribution-NonCommercial 2.5 License.
- Students working together. By Yuuki Guzman and Agoston Tyll, Okinawa Institute of Science and Technology, 2015.
- Two pints of Guinness. By Harald Bischoff - Own work, CC BY-SA 3.0,
<https://commons.wikimedia.org/w/index.php?curid=23344636>